


Diffusion-LM Improves Controllable Text Generation

Xiang Lisa Li, John Thickstun, Ishaan Gulrajani, Percy Liang, Tatsunori B. Hashimoto
Stanford University



Motivating Diffusion-LM




(Dhariwal and Nichol, 2021)

Diffusion models are now dominant in vision. Are they also good for language?


Motivating Diffusion-LM: Classifier Guidance

- Diffusion models can be easily and convincingly steered using a probabilistic scoring function (e.g., a classifier).
- Analogous to “plug-and-play” language modeling (Dathathri et al., 2020).
- This post-hoc conditioning seems compelling:
 - train one general-purpose (expensive) generative model.
 - steer it for your specialized task at inference time.



(Dhariwal and Nichol, 2021)

Motivating Diffusion-LM



vibrant portrait painting of Salvador Dalí with a robotic half face



a shiba inu wearing a beret and black turtleneck
(Ramesh et al., 2022)



a close up of a handpalm with leaves growing from it

Classifier Guidance vs. Prompting: competing or complementary paradigms?

Diffusion-LM

Introduction

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Control
- Experiments
- Discussion


Talk Outline

- **Constructing Diffusion-LM**
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Control
- Experiments
- Discussion

Talk Outline


- **Constructing Diffusion-LM**
 - ★ **The Standard Diffusion Model (Ho et. al., 2020)**
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Control
- Experiments
- Discussion

Denoising Diffusion Probabilistic Models



- Learn to generate data by progressive denoising.

Denoising Diffusion Probabilistic Models




- **Forward Process.** Given *noise schedule* β_1, \dots, β_T (hyper-parameters):

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N} \left(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t I \right).$$

- **Reverse Process.** Learn to denoise with parameters θ :

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N} \left(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t) \right).$$

Optimizing a Diffusion Model




- Optimize like a VAE (the usual variational lower bound):

$$\begin{aligned}
 -\log p_\theta(\mathbf{x}_0) &= -\log \int_{\mathbf{x}_{1:T}} p_\theta(\mathbf{x}_{0:T}) d\mathbf{x}_{1:T} = -\log \mathbb{E}_{\mathbf{x}_{1:T} \sim q} \left[\frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] \\
 &\leq \mathbb{E}_{\mathbf{x}_{1:T} \sim q} \left[-\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] = \mathbb{E}_{\mathbf{x}_{1:T} \sim q} \left[-\log q(\mathbf{x}_T) - \sum_{t=1}^T \log \frac{p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q(\mathbf{x}_t | \mathbf{x}_{t-1})} \right].
 \end{aligned}$$

- Many small steps β_t make the Gaussian approximation $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ valid.

In What Sense is this Denoising?



- Rewrite the variational objective (algebra; see Ho et al., 2020; Appendix A):

$$\mathbb{E}_{\mathbf{x}_{1:T} \sim q} \left[D(q(\mathbf{x}_T | \mathbf{x}_0) \| p(\mathbf{x}_T)) + \sum_{t=1}^T D(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) \| p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)) - \log p_\theta(\mathbf{x}_0 | \mathbf{x}_1) \right].$$

- Want to make $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ look like the posterior distribution $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$.


Talk Outline

- **Constructing Diffusion-LM**
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Talk Outline

- **Constructing Diffusion-LM**
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ **Learning Word Embeddings (End-to-End Training)**
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Learning Word Embeddings




- Standard DDPM assumes inputs \mathbf{x}_0 are continuous.
- Language model inputs \mathbf{w} are discrete tokens.
- Use the re-parameterization trick to learn word embeddings $q_\phi(\mathbf{x}_t | \mathbf{w})$:

$$-\log p_\theta(\mathbf{x}_0) \leq \mathbb{E}_{\substack{\mathbf{x}_{1:T} \sim q \\ \mathbf{x}_0 \sim q_\phi}} \left[-\log q(\mathbf{x}_T) - \sum_{t=1}^T \log \frac{p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)}{q(\mathbf{x}_t | \mathbf{x}_{t-1})} - \log p_\theta(\mathbf{w} | \mathbf{x}_0) \right].$$

Are These Learned Embeddings Meaningful?

- Learning the embedding seems important.
- Random embeddings performed poorly.
- Pre-trained embeddings from an AR model also performed poorly.



A t-SNE plot of learned embeddings, colored according to POS.


Talk Outline

- **Constructing Diffusion-LM**
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Talk Outline

- **Constructing Diffusion-LM**
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ **Predicting the Noiseless Embeddings**
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Predicting the Noiseless Embeddings



- We want to minimize terms $D(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) \parallel p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t))$, where
$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t)).$$
- Closed form for the posteriors: $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t I)$, where $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) = r_t \mathbf{x}_t + s_t \mathbf{x}_0$.
- And $r_t, s_t, \tilde{\beta}_t$ are constants derived from the noise schedule β_1, \dots, β_T .

Predicting the Noiseless Embeddings

- Directly parameterize $\mu_\theta(\mathbf{x}_t, t)$ to approximate $\tilde{\mu}_t(\mathbf{x}_t, \mathbf{x}_0) = r_t \mathbf{x}_t + s_t \mathbf{x}_0$?
- But we already know \mathbf{x}_t .
- Ho et al., 2020: write $\mathbf{x}_0 = \mathbf{x}_t - \varepsilon_t$, predict $\varepsilon_t \approx \varepsilon_\theta(\mathbf{x}_t, t)$ and reparameterize

$$\mu_\theta(\mathbf{x}_t, t) = (r_t + s_t)\mathbf{x}_t - s_t \varepsilon_\theta(\mathbf{x}_t, t).$$

- Li et al., 2022: predict $\mathbf{x}_0 \approx f_\theta(\mathbf{x}_t, t)$ and reparameterize

$$\mu_\theta(\mathbf{x}_t, t) = r_t \mathbf{x}_t + s_t f_\theta(\mathbf{x}_t, t).$$


Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- **Sampling from Diffusion-LM**
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Talk Outline


- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Sampling from Diffusion-LM




- Sample from a Gaussian $\mathbf{x}_T \sim \mathcal{N}(0, I)$ and then just follow the chain.
- Sampling Diffusion-LM requires T (diffusion steps) model calls.
- In contrast, sampling AR models requires L (sequence length) model calls.
- In our experiments, $T \gg L$ so sampling is slow.

Sampling Heuristics



- Analogous AR sampling heuristics can be applied to Diffusion-LM.
- Temperature sampling: reduce the noise at each sampling step.
- Nucleus sampling: truncate the tails of the Gaussian noise.
- Minimum Bayes Risk (MBR) decoding.

The Clamping Trick



- At each step we predict $\mathbf{x}_0 \approx f_\theta(\mathbf{x}_t, t)$.
- *The clamping trick:* instead predict the nearest embedding in the dictionary:
$$\mathbf{x}_0 \approx \text{Clamp}(f_\theta(\mathbf{x}_t, t)).$$
- This seems to nudge Diffusion-LM to commit to tokens earlier.

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- **Sampling from Diffusion-LM**
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- **Sampling from Diffusion-LM**
 - ★ The Clamping Trick and Other Heuristics
 - ★ **Classifier-Guided Sampling**
- Experiments
- Discussion

Classifier-Guided Sampling

- Given: labeled data pairs (\mathbf{x}, \mathbf{c}) describing the desired attribute \mathbf{c} .
- Goal: sample from the posterior distribution $p_\theta(\mathbf{x}_{0:T}|\mathbf{c}) = \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{c})$.
- Each term can be rewritten as $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{c}) \propto p_\theta(\mathbf{c}|\mathbf{x}_{t-1}, \mathbf{x}_t)p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$.
- The term $p_\theta(\mathbf{c}|\mathbf{x}_{t-1}, \mathbf{x}_t)$ is a classifier. And $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$ is Diffusion-LM.
- In practice, the classifier doesn't need to see both \mathbf{x}_t and \mathbf{x}_{t-1} .
- Train the classifier $p_\theta(\mathbf{c}|\mathbf{x}_{t-1})$ on noisy data.

Langevin Dynamics

- Goal: sample from the posterior $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{c})$.
- Langevin Dynamics: define a Markov chain

$$\mathbf{x}_{t-1}^{(i+1)} = x_{t-1}^{(i)} - \eta \nabla_{\mathbf{x}_{t-1}} \log p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{c}) + \sqrt{2\eta} \varepsilon_i, \text{ where } \varepsilon_i \sim \mathcal{N}(0, I).$$


- For small η , as $i \rightarrow \infty$, $D(\mathbf{x}_{t-1}^{(i+1)} \parallel p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{c})) \rightarrow 0$. The gradient is:

$$\nabla_{\mathbf{x}_{t-1}} \log p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{c}) = \boxed{\nabla_{\mathbf{x}_{t-1}} \log p_\theta(\mathbf{c}|\mathbf{x}_{t-1})} + \boxed{\nabla_{\mathbf{x}_{t-1}} \log p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)}.$$

Classifier Score Diffusion-LM

- In practice, init $\mathbf{x}_{t-1}^{(0)} \sim p_\theta(\mathbf{x}_{t-1}, \mathbf{x}_t^{(1)})$ (warmstart) and take just one step.

Iterative Gradient-Based Control



- Conceptually similar to PPLM (Dathathri et al., 2020).
- But control is applied coarse-to-fine, instead of left-to-right.

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- **Experiments**
- Discussion

Datasets

- Two datasets: **E2E** and **ROCStories**.
- **E2E**. 50k restaurant reviews. Sample text: “Browns Cambridge is good for Japanese food and also children friendly near The Sorrento.”
- **ROCStories**. 98k short stories. Sample text: "Jennifer has a big exam tomorrow. She got so stressed, she pulled an all-nighter. She went into class the next day, weary as can be. Her teacher stated that the test is postponed for next week. Jennifer felt bittersweet about it."
- Small datasets: scaling up Diffusion-LM is an open problem.

Control Tasks

input (Semantic Content)	food : Japanese
output text	Browns Cambridge is good for Japanese food and also children friendly near The Sorrento .
input (Parts-of-speech)	PROPN AUX DET ADJ NOUN NOUN VERB ADP DET NOUN ADP DET NOUN PUNCT
output text	Zizzi is a local coffee shop located on the outskirts of the city .
input (Syntax Tree)	(TOP (S (NP (*) (*) (*)) (VP (*) (NP (NP (*) (*)))))))
output text	The Twenty Two has great food
input (Syntax Spans)	(7, 10, VP)
output text	Wildwood pub serves multicultural dishes and is ranked 3 stars
input (Length)	14
output text	Browns Cambridge offers Japanese food located near The Sorrento in the city centre .
input (left context)	My dog loved tennis balls.
input (right context)	My dog had stolen every one and put it under there.
output text	One day, I found all of my lost tennis balls underneath the bed.

- Six controllable generation tasks.

Baselines

- For classifier-guided control:
 - PPLM (Dathathri et al., 2020).
 - FUDGE (Yang and Klein, 2021).
 - Finetuning (skyline).
- For infilling:
 - DELOREAN (Qin et al., 2020).
 - COLD (Qin et al., 2021).
 - Finetuning (skyline).

A Qualitative Example

Syntactic Parse	(S (S (NP *) (VP * (NP (NP * *) (VP * (NP (ADJP * *) *))))) * (S (NP * * *) (VP * (ADJP (ADJP *)))))
FUDGE	Zizzi is a cheap restaurant . [incomplete]
Diffusion-LM	Zizzi is a pub providing family friendly Indian food Its customer rating is low
FT	Cocum is a Pub serving moderately priced meals and the customer rating is high
Syntactic Parse	(S (S (VP * (PP * (NP * *)))) * (NP * * *) (VP * (NP (NP * *) (SBAR (WHNP *) (S (VP * (NP * *)))))) *)
FUDGE	In the city near The Portland Arms is a coffee and fast food place named The Cricketers which is not family - friendly with a customer rating of 5 out of 5 .
Diffusion-LM	Located on the riverside , The Rice Boat is a restaurant that serves Indian food .
FT	Located near The Sorrento, The Mill is a pub that serves Indian cuisine.

- FUDGE and Finetuning (FT) deviate after a few tokens (exposure bias).
- Diffusion-LM is robust to local failures to apply the control.

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Talk Outline

- Constructing Diffusion-LM
 - ★ The Standard Diffusion Model (Ho et. al., 2020)
 - ★ Learning Word Embeddings (End-to-End Training)
 - ★ Predicting the Noiseless Embeddings
- Sampling from Diffusion-LM
 - ★ The Clamping Trick and Other Heuristics
 - ★ Classifier-Guided Sampling
- Experiments
- Discussion

Limitations

- Decoding is much slower than AR Transformer models: 2000 diffusion steps versus (short) sequence-length AR steps.
- Training seems to converge more slowly than for AR models.
- Diffusion-LM has higher perplexity than comparably-sized AR models.
- Similar challenges for diffusion models of images have been observed and overcome, so there is reason to be hopeful!

Coarse-to-Fine Generation

- Coarse-to-fine generation seems helpful for control (versus autoregressive generation) because the control target can be incorporated globally into the plan for generating text.
- You can use Langevin dynamics to sample coarse-to-fine from AR models fine-tuned in noisy data (Jayaram and Thickstun, 2021)
- A mystery: I spent several months trying to adapt these methods to text, but couldn't get the sampling to work very well.
- Maybe coarse-to-fine isn't the only thing that's going right here?

Non-Autoregressive Language Generation

- Previously non-autoregressive open-ended language generation has seemed difficult (GAN, VAE).
- Adapting the DDPM recipe to text required some alterations, but if similarly-scoped alterations would make VAE's work well for text it seems like this would have happened by now.
- What (if anything) is different about diffusion models?

Continuous vs. Discrete Models

- NLP folks spend a lot of time trying to convert their discrete data into continuous representations.
- Vision folks spend a lot of time trying to convert their continuous data into discrete representations.
 - VQ-VAE2 (Razavi et al., 2019)
 - VQ-GAN (Esser et al., 2021)
 - Parti (Yu et al., 2022)
- What is going on here?

Thank You!

Paper: <https://arxiv.org/abs/2205.14217>

Code: <https://github.com/XiangLi1999/Diffusion-LM>

Diffusion-LM